

Field Notes

**Master path for AI-assisted knowledge work —
one git exemplar**

Henrik Hellbusch · 2026

How this deck is organized

- **Foundations** — harness thesis, context, mechanism, epistemics, the shift
- **Memory & State** — chat limits, artifacts, memory taxonomy
- **Knowledge Systems** — PKM, pipeline, honest limits
- **Practice** — verify, spar, shoshin, the long game
- **The Workspace** — repo map, toolkit, proactive runtime

Appendix = reading map. Through-line: trust, but verify.

Foundations

Everything is an LLM call

All agent systems → calls to a model provider.

Only difference: the **context** bundled each call — the harness.

Talk: `alex-krentsel-openclaw-deep-dive.md` (Krentsel / OpenClaw architecture)

The harness may matter more than the model

Same model, different harness → very different outcomes.

Terminal Bench: Claude Code ~#40 · same Opus in other harnesses ~#1.

When models are capable enough, the bottleneck is **reliable operation** — not raw IQ.

Library: `alberta-tech-why-devs-obsessed-claude-code.md` ·
`ryan-lopopolo-harness-engineering.md`

Prompt · context · harness

Layer	Question
Prompt	What to ask?
Context	What to send?
Harness	How does the system run?

Tools, permissions, tests, retries, guardrails — not just tokens in the window.

Library: `tejas-kumar-harnesses-in-ai.md` · Lopopolo · Horthy

The harness

Each model call receives a **bundle of context**:

- Your messages · prior turns · files in scope
- Rules and skills · results from tools

Harness = how that bundle is assembled each turn.

The harness — why structure wins

Clear folders and explicit rules beat a clever one-liner.

Same for code, notes, or a pile of PDFs.

Harness building blocks

Piece	When loaded	Field Notes
AGENTS.md	Every session	Workspace contract
Rules	Always-on	<code>.cursor/rules/</code>
Skills	When invoked	<code>.agents/skills/</code>

Agent Skills — agentskills.io

- **Skill** = workflow recipe — not an MCP server
- Progressive load: name → `SKILL.md` → linked files
- `/start`, `/spar`, `/review` — portable across Cursor, Pi, Claude Code

Guide: `framework-bootstrap.md`

The context window

- Fixed **token budget** for every session
- Shared by: messages, history, loaded files, tools, rules, skills
- Not infinite — and not all of it is usable

Design question: what earns a place in that budget?

The dumb zone

- Past ~40% **context fill**, model reliability drops
- Verbose tool output, many MCPs, long correction chains — all push you there
- **Design for the smart zone:** fresh sessions, intentional scope, RPI

Talk: `dex-horthy-no-vibes-allowed.md`

Context budget — what draws from it

Item	Loaded when
AGENTS.md	Every session
Rules	Always-on
Skills	When invoked
Open files / tools	Harness decides

Every addition is a **tradeoff** — budget it deliberately.

The KV cache — why context has a ceiling

- Context window = a **fixed-size allocation** at session start — not a sliding window
- Allocated from VRAM remaining after weights load → runtime size varies
- **Can't resize mid-session** — earlier tokens evict when it fills

Skip ahead if you don't run local models — the dumb zone is what matters in practice.

Deeper: `docs/ai-engineering/what-a-context-window-actually-is.md`

Software 3.0

Context window = the programming surface.

`AGENTS.md` · rules · skills → not documentation — **they are the code.**

Quality of context determines quality of output.

Talk: `andrej-karpathy-vibe-coding-to-agentic-engineering.md`

Harness problems — the interface layer

- Harness $\neq$ model only — **editor and product** assemble the bundle
- **Greedy context:** assumes maximum help; fills the window fast
- **Editor bleed:** files in prompt because they were *open*, not in scope
- **Session blend:** unrelated tabs and old chats mixed into one thread

Mitigate: fresh sessions · explicit scope · rules for in/out

Under the hood — the loop stack

What the harness wraps.

Skip to Act II if you just want the practices — come back when model behavior surprises you.

Transformer inference

- **Transformer** — architecture; **weights** learned during training
- **Inference** — forward pass at runtime (**not** training)
- **One step**: tokens in $\rightarrow$ **$P(\text{next token} \mid \text{so far})$** $\rightarrow$ pick one $\rightarrow$ append

One call $\rightarrow$ **one token** — not a whole answer at once.

Transformer inference — the loop

- Append the chosen token → feed the string back → infer again
- A sentence, paragraph, or story = **many** such steps
- **Loop 1** (next slide) is this repetition — Krentsel's first "loopiness"

Sampling (temperature, top-p) usually sits in the **LLM/harness** layer.

Loop 1 — next token

```
tokens in → Transformer → next token → append → ...
```

The  is the loop above — one token per pass.

Loop 2 — chat

```
You → LLM → reply
  ↑_____| (conversation)
```

Assistant phase — reactive, turn by turn.

Loop 3 — scoped agent

```
Goal → LLM → tool → result
      ↑___observe___|   (closed control loop)
```

Tools + **act** → **see outcome** → **decide next** — not one-shot Q&A.

Same pattern for code, research, and docs.

Loop 4 — autonomous (outer shell)

Broader scope: dynamic tools, self-config, cron/heartbeat.

OpenClaw-class runtimes sit here; Cursor/Pi often sit at Loop 3.

Each layer **wraps** the loops inside — matryoshka.

The nested stack

↳ Transformer Inference

Large-Language Model

Assistants – ChatGPT, Claude, Gemini

Scoped agents + tooling – Claude Code, Codex, Cursor

Autonomous + env. ownership – OpenClaw

Each line **wraps** the one above — matryoshka (Krentsel).

Nested stack — inner two layers

Layer	What one "unit" does
Transformer inference	One forward pass → one next token
Large-Language Model	Many steps — tokenizer, API, sampling

Chat, agents, and harnesses **repeat** these inner calls with more context.

Talk: `alex-krentsel-openclaw-deep-dive.md` (~5:13–7:09)

Statistical, not deterministic

- **Traditional code** — logic you wrote; same input → same output
- **LLM** — statistical model; $P(\text{next token} \mid \text{context})$ each step
- **Sampling** — same prompt can yield different answers (*at non-zero temperature*)

Not broken — **plausible**, not guaranteed correct.

Why that changes how you work

- Predicts what **reads well** — not what **is** true
- Mistakes look confident — no crash, no stack trace
- **Harness + verify** — tests, lint, re-read sources

Act IV: **trust, but verify** — same thread.

The shift — what got cheaper

For decades, **implementation** was often the bottleneck.

AI **compresses** it — does not eliminate verification.

Essay: `the-shift.md`

The shift — what matters more

When implementation is cheap, value **often** moves to:

- **Decomposition** — what are we actually building?
- **Verification** — is this correct, safe, sourced?
- **Judgment** — context, ethics, what you sign your name to

Jagged intelligence

Models peak where **RL can verify** — code, math, structured tasks.

Weak where verification is hard — novel judgment, embodied reasoning, ethics.

"What stays human" isn't vague — it's the unverifiable gap.

Talk: `andrej-karpathy-vibe-coding-to-agentic-engineering.md`

Fluency isn't evidence

- AI output looks authoritative — same confidence, right or wrong
- Data-driven posture: *show me the evidence* applies here too
- Treat fluent output like an **unvalidated claim**

Trust, but verify

Trust the speed — drafts and first passes are useful.

Verify before you rely — test, source-check, second look.

Fluent $\neq$ correct — code, citations, advice.

Sycophancy and agreement bias

- "Great idea." "You're right." — praise without evaluation
- Mirrors your frame — confident drafts that *sound* like you
- **Sycophancy**: agreeableness structurally rewarded (RLHF)
- Root: **frictionless** — feels complete without challenge

Essay: `the-shift.md` §6

Push back on agreeableness

- **Treat agreement as a null signal** — validates almost everything
- **No pleasantries** — rule or prompt: skip praise and filler
- **Ask what's wrong** — spar, adversarial review, human pushback

Case study: `frictionless-entity.md`

AI amplifies how you already work

- Tools **amplify** habits and culture — good and bad
- Strong review, handoffs, verification → **compounds faster**
- Shortcuts, rubber-stamping → **compounds faster too**

AI adoption — not bolt-on

Real adoption means **re-evaluating how you work** — not adding a chat box.

Essay: `the-shift.md` — culture amplification

What stays human

When AI drafts faster, you still bring:

- **Clarity** — what is the question? What would change your mind?
- **Verification** — before you stand behind output
- **Judgment** — trade-offs, audience, professional standards
- **Memory** — write down what survived verification

Memory & State

The chat problem — cross-session

New session = blank slate. Prior decisions aren't loaded unless you put them there.

Long chats **compress** earlier context — details become approximations.

Essay: `prompting-and-state.md`

The chat problem — what we skip

"I'll write it up later" rarely happens.

The thread dies; the learning evaporates.

Saved artifacts are memory

“ **Files you keep** are memory. **Chat is not.** ”

Findable · revisable · linkable · diffable.

Saved artifacts — git when ready

A folder works to start.

Git when you want history, branches, and shared review.

Artifacts as async multiplier

- **Cognitive offload** — write it → safely flush working memory → recall on demand
- **Async self-service** — peer reads the doc; you're not the bottleneck for first contact
- **Teaching economy** — your time goes to judgment and troubleshooting, not procedure transfer

The doc handles the rote. You handle what can't be documented.

Two kinds of memory (Simon Scrapes)

Kind	Question	Field Notes
Operational (L1–4)	What did we decide?	Strong — manual L1–2
Knowledge (L5–6)	What did I learn?	Strong on L5 . library/

Don't conflate decision memory with curated learning.

Library: `simon-scrapes-claude-code-memory-systems.md`

Operational memory — L1 and L2

Level 1: always-loaded context (`AGENTS.md`, brief)

Level 2: handoffs — backlog, `whats-next`, session rules

Operational memory — L3+

Semantic search, cron, heartbeat — when *find-it-again* or *autonomous agents* hurt.

Start manual. Automate when friction is real.

Knowledge Systems

Three conversations — usually separate

Conversation	Examples	Field Notes analogue
PKM / second brain	PARA, Zettelkasten, digital garden	capture → distill → express
LLM knowledge base	Karpathy LLM Wiki, Recall	research/ → library/
Agent memory / harness	hooks, memsearch, OpenClaw	rules, skills, handoffs

Three conversations — one git map

This repo **maps all three in git** — one implementation, not the only layout.

Map the pipeline to your stack

Pattern: raw capture → curated wiki → finished output

Paths: `research/` → `library/` → `docs/` · `devops/`

Map the pipeline — schema

Schema matters more than folder names — what goes where, how it's reviewed.

Where this exemplar stops — strong

Strong today: manual L1–2 · L5 (library/)

Rules, handoffs, backlog · curated library wiki.

Where this exemplar stops — gaps

Not default: automated L3–4 · L6 (shared agent brain)

Add when *find-it-again* or *autonomous agents* hurt — not day one.

Capture → synthesize → express

```
research/ → library/ → docs/ · devops/  
(raw)      (wiki)      (finished)
```

Research = drawer · **Library** = curated wiki · **Docs** = field notes for others.

Capture pipeline — detail

- **Research** — transcripts, refs, experiment logs
- **Library** — summaries, themes, links
- **Docs / devops** — essays and runnable examples

Karpathy's three layers

LLM Wiki	Field Notes
<code>raw-sources/</code>	<code>research/*/sources/</code>
<code>wiki/</code>	<code>library/</code>
<code>schema/</code>	<code>AGENTS.md</code> , <code>ingest</code> , <code>/audit</code>

Karpathy — why schema matters

Raw pile $\neq$ knowledge base.

Schema — where things go, how they're reviewed — is the third layer.

Library: `karpathy-11m-wiki.md`

Zettelkasten vs essay grain

Zettelkasten: atomic notes, stable IDs, links emerge bottom-up.

Field Notes: synthesis grain — essays, case studies, runnable examples.

Both valid. Essay/case-study grain here — atomic Zettel optional.

When to add a database

Don't replace markdown with SQL by default.

Add vector/graph/Postgres when *find-it-again* or *multi-agent recall* hurts more than curation cost.

Practice

Your first session — three paths

A — code: small task · review diff · test · commit

B — learning: one source · structured notes · **save to a file**

C — chat only: one question — won't compound without B

First session — shared rules

Trust, but verify · small scope · no secrets in prompts

Case study: `docs/ai-engineering/ai-for-unfamiliar-domains.md`

Session habits that compound

Habit	Defends against
Backlog / inbox	"I'll remember later"
Handoff when pausing	context loss between sessions
Re-read source files	in-session compaction

Essay: `session-framework.md`

Session habits (continued)

Habit	Defends against
Checkpoint commits	crash mid-session
Review before merge	fluent wrong output shipping

Add **rules and skills** when friction repeats — not before.

The long game

After 6 months of session discipline:

- **library/** — curated sources you can cite, not chat logs you can't find
- **Rules and skills** — compound; each one makes future sessions faster
- **Review trail** — you know what you verified, when, at which SHA

A pure-chat user has nothing recoverable. You have a workspace.

Sparring and shoshin — the bracket

shoshin → draft → spar → revise → share

Sparring and shoshin — when

Shoshin (start): read sources; not stale summaries.

Spar (after draft): steel-man counterarguments — thesis, not typos.

Essay: `docs/ai-engineering/sparring-and-shoshin.md`

Zanshin — remaining mind

Zanshin (残心): mind that **remains through** the technique.

The model carries nothing. Framework practices restore what distraction breaks.

Essay: `zanshin.md`

Learning capacity — the full cup

When people are **too full to absorb**, more input spills.

Teams adopting AI: draft throughput can exceed verify capacity.

Essays: `the-full-cup.md` · practitioners guide

The Workspace

What Field Notes is

Public **git notebook** — built from real work over time.

Not a product. Essays · curated library · philosophy · optional infra reference · session rules.

Browse selectively. Fork ideas, ignore rooms that aren't yours.

What's in the repo — core

- `docs/` — essays and case studies
- `library/` — curated sources
- `research/` — raw capture → library
- `devops/` — infra examples (*optional*)

Index: `docs/README.md`

What's in the repo — working layer

- `AGENTS.md` — workspace contract (always loaded)
- `.agents/skills/` · `.cursor/rules/` — skills + rules
- `BACKLOG.md` · `.planning/` — projects in flight

Four ways one notebook serves different readers

1. **Learning** — sources, summaries, experiment logs
2. **Field notes** — write-ups you reuse in other contexts
3. **What people ask for** — links, examples, short guides
4. **Projects in flight** — backlog, handoffs, threads

Four readers — one notebook

Future you · peers · collaborators · tools — different rooms, same artifacts.

Portable toolkit — three layers

Layer	Solves	This workspace
Runtime	where the agent runs; git sync	Paude + Pi
Discipline	verify, spar, session tracking	Zanshin kit
Workspace	accumulated context	this repo

Portable toolkit — one command

Mount runtime + discipline + workspace on a new problem.

Essay: `portable-ai-toolkit.md` (*working draft*)

Brain + proactive runtime

Brain — git knowledge you own: `library/`, `docs/`, rules, `/audit`.

Proactive runtime (*OpenClaw-class*) — cron, heartbeat, channels, shell.

Complementary — not either/or.

Brain + runtime — how they connect

Field Notes (git)	OpenClaw-class runtime
library/, docs/, rules	cron, heartbeat, channels
└── workspace mounted ───┘	

Promotion rule: commit to `library/` or `BACKLOG` — not only vector memory.

Library: `alex-krentsel-openclaw-deep-dive.md` (*exploratory*)

What this deck does *not* claim

- One size fits all — **browse selectively**
- Replace your judgment or professional standards

What this deck does *not* claim (continued)

- All repo content is author-reviewed (*much is draft*)
- You need OpenClaw, Paude, or this exact layout

Trust, but verify. The path is yours.

Thanks

Start anywhere — each section stands alone.

1. One place for things worth finding again
2. One bounded AI task — **trust, but verify**
3. One link from the **appendix**

Appendix

Reading map — what to read and why

Pick **one** row. Trust, but verify.

Start here

`the-shift.md`

Why the bottleneck moved — and why AI amplifies the culture you already have.

`ego-ai-and-the-zen-antidote.md`

Why AI can reinforce bad framing — and what to do.

Start here — state

`prompting-and-state.md`

Why a great chat fails across days or weeks.

Start here (continued)

`zanshin.md`

What should survive when the session ends.

`docs/ai-engineering/sparring-and-shoshin.md`

Adversarial review and beginner's mind — no tooling required.

Start here — verification

`frictionless-entity.md`

Root property behind agreement bias and inherited framing.

`docs/ai-engineering/ai-for-unfamiliar-domains.md`

Verification in practice: right answer, wrong first draft.

Learning & teams

`the-full-cup.md`

When people are too full to learn.

`the-full-cup-practitioners-guide.md`

Companion playbook if the essay landed.

Learning & teams (continued)

`session-framework.md`

Session habits — what each defends against and why.

Workflows

`ai-assisted-development-workflows.md`

Multi-session patterns + per-tool context files.

`framework-bootstrap.md`

Load Zanshin posture + AGENTS.md / skills / rules by tool.

Workflows (continued)

`docs/case-studies/building-knowledge-management-with-ai.md`

How backlog, library, and session tools were built.

Memory & PKM

`library/README.md`

Curated talks and articles — learn without starting from zero.

`simon-scrapes-claude-code-memory-systems.md`

Six-level agent memory taxonomy — comparison vocabulary.

Memory & PKM (continued)

`karpathy-llm-wiki.md`

Raw sources vs wiki — why a pile of notes isn't a knowledge base.

`alex-krentsel-openclaw-deep-dive.md`

Agentic loop matryoshka — harness, gateway, proactive runtime.

Harness engineering (AI Engineer)

`ryan-lopopolo-harness-engineering.md`

Harness is the job — docs, lint, reviewer agents; code as build artifact.

`dex-horthy-no-vibes-allowed.md`

Context engineering — stay out of the "dumb zone"; RPI workflow.

Harness engineering (continued)

`tejas-kumar-harnesses-in-ai.md`

First-principles anatomy + demo — verify step; prompt unchanged.

`alberta-tech-why-devs-obsessed-claude-code.md`

Same model, different harness — Terminal Bench ranks the multiplier.

Philosophy & human layer

`docs/case-studies/who-is-speaking.md`

AI writes in your voice — biographical claims, professional identity, opinions attributed to you. Adjacent to sycophancy; a distinct failure mode.

`docs/philosophy/the-dojo-after-the-automation.md`

Who builds the humans capable of directing AI? The capacity to judge and set direction doesn't emerge automatically when execution automates.

Index & toolkit

`portable-ai-toolkit.md`

Runtime + discipline + workspace (*working draft*).

`docs/README.md` · `devops/` (*optional*)

Master index · runnable infra.